

# credit\_risk\_ensemble

May 30, 2021

## 1 Ensemble Learning

### 1.1 Initial Imports

```
[1]: import warnings
warnings.filterwarnings('ignore')
```

```
[2]: import numpy as np
import pandas as pd
from pathlib import Path
from collections import Counter
from datetime import datetime
```

```
[34]: from sklearn.metrics import balanced_accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.datasets import make_classification

from imblearn.metrics import classification_report_imbalanced
from imblearn.ensemble import BalancedRandomForestClassifier
from imblearn.ensemble import EasyEnsembleClassifier
```

### 1.2 Read the CSV and Perform Basic Data Cleaning

```
[4]: # Load the data
file_path = Path('Resources/LoanStats_2019Q1.csv')
df = pd.read_csv(file_path)

# Preview the data
df.head()
```

```
[4]:   loan_amnt  int_rate  installment  home_ownership  annual_inc  \
0    10500.0    0.1719         375.35             RENT      66000.0
1    25000.0    0.2000         929.09             MORTGAGE    105000.0
```

2	20000.0	0.2000	529.88	MORTGAGE	56000.0
3	10000.0	0.1640	353.55	RENT	92000.0
4	22000.0	0.1474	520.39	MORTGAGE	52000.0

	verification_status	issue_d	loan_status	pymnt_plan	dti	...	\
0	Source Verified	Mar-2019	low_risk	n	27.24	...	
1	Verified	Mar-2019	low_risk	n	20.23	...	
2	Verified	Mar-2019	low_risk	n	24.26	...	
3	Verified	Mar-2019	low_risk	n	31.44	...	
4	Not Verified	Mar-2019	low_risk	n	18.76	...	

	pct_tl_nvr_dlq	percent_bc_gt_75	pub_rec_bankruptcies	tax_liens	\
0	85.7	100.0	0.0	0.0	
1	91.2	50.0	1.0	0.0	
2	66.7	50.0	0.0	0.0	
3	100.0	50.0	1.0	0.0	
4	100.0	0.0	0.0	0.0	

	tot_hi_cred_lim	total_bal_ex_mort	total_bc_limit	\
0	65687.0	38199.0	2000.0	
1	271427.0	60641.0	41200.0	
2	60644.0	45684.0	7500.0	
3	99506.0	68784.0	19700.0	
4	219750.0	25919.0	27600.0	

	total_il_high_credit_limit	hardship_flag	debt_settlement_flag
0	61987.0	N	N
1	49197.0	N	N
2	43144.0	N	N
3	76506.0	N	N
4	20000.0	N	N

[5 rows x 86 columns]

```
[5]: df.columns
```

```
[5]: Index(['loan_amnt', 'int_rate', 'installment', 'home_ownership', 'annual_inc',
        'verification_status', 'issue_d', 'loan_status', 'pymnt_plan', 'dti',
        'delinq_2yrs', 'inq_last_6mths', 'open_acc', 'pub_rec', 'revol_bal',
        'total_acc', 'initial_list_status', 'out_prncp', 'out_prncp_inv',
        'total_pymnt', 'total_pymnt_inv', 'total_rec_prncp', 'total_rec_int',
        'total_rec_late_fee', 'recoveries', 'collection_recovery_fee',
        'last_pymnt_amnt', 'next_pymnt_d', 'collections_12_mths_ex_med',
        'policy_code', 'application_type', 'acc_now_delinq', 'tot_coll_amt',
        'tot_cur_bal', 'open_acc_6m', 'open_act_il', 'open_il_12m',
        'open_il_24m', 'mths_since_rcnt_il', 'total_bal_il', 'il_util',
        'open_rv_12m', 'open_rv_24m', 'max_bal_bc', 'all_util',
```

```

'total_rev_hi_lim', 'inq-fi', 'total_cu_tl', 'inq_last_12m',
'acc_open_past_24mths', 'avg_cur_bal', 'bc_open_to_buy', 'bc_util',
'chargeoff_within_12_mths', 'delinq_amnt', 'mo_sin_old_il_acct',
'mo_sin_old_rev_tl_op', 'mo_sin_rcnt_rev_tl_op', 'mo_sin_rcnt_tl',
'mort_acc', 'mths_since_recent_bc', 'mths_since_recent_inq',
'num_accts_ever_120_pd', 'num_actv_bc_tl', 'num_actv_rev_tl',
'num_bc_sats', 'num_bc_tl', 'num_il_tl', 'num_op_rev_tl',
'num_rev_accts', 'num_rev_tl_bal_gt_0', 'num_sats', 'num_tl_120dpd_2m',
'num_tl_30dpd', 'num_tl_90g_dpd_24m', 'num_tl_op_past_12m',
'pct_tl_nvr_dlq', 'percent_bc_gt_75', 'pub_rec_bankruptcies',
'tax_liens', 'tot_hi_cred_lim', 'total_bal_ex_mort', 'total_bc_limit',
'total_il_high_credit_limit', 'hardship_flag', 'debt_settlement_flag'],
dtype='object')

```

```

[6]: ## Adding month and date columns from dates

# df["month"] = datetime.strptime(df["issue_d"], 'MMM-YYYY').month
# >>> print dt.year, dt.month, dt.day
df['issue_date'] = pd.to_datetime(df['issue_d'])
df['next_pymnt_date'] = pd.to_datetime(df['next_pymnt_d'])

df['issue_month'] = pd.DatetimeIndex(df['issue_date']).month.astype(str)
df["issue_year"] = pd.DatetimeIndex(df['issue_date']).year.astype(str)
df['next_pymnt_d_month'] = pd.DatetimeIndex(df['next_pymnt_date']).month.
    ↳astype(str)
df['next_pymnt_d_year'] = pd.DatetimeIndex(df['next_pymnt_date']).year.
    ↳astype(str)

```

```

[7]: df

```

```

[7]:      loan_amnt  int_rate  installment  home_ownership  annual_inc  \
0      10500.0    0.1719      375.35          RENT      66000.0
1      25000.0    0.2000      929.09      MORTGAGE    105000.0
2      20000.0    0.2000      529.88      MORTGAGE     56000.0
3      10000.0    0.1640      353.55          RENT     92000.0
4      22000.0    0.1474      520.39      MORTGAGE     52000.0
...      ...      ...      ...      ...      ...
68812    10000.0    0.1502      346.76          RENT     26000.0
68813    12000.0    0.2727      368.37          RENT     63000.0
68814     5000.0    0.1992      185.62      MORTGAGE     52000.0
68815    40000.0    0.0646     1225.24      MORTGAGE    520000.0
68816    16000.0    0.1131      350.36      MORTGAGE     72000.0

      verification_status  issue_d  loan_status  pymnt_plan    dti  ...  \
0          Source Verified  Mar-2019    low_risk          n  27.24  ...
1              Verified  Mar-2019    low_risk          n  20.23  ...
2              Verified  Mar-2019    low_risk          n  24.26  ...

```

3	Verified	Mar-2019	low_risk	n	31.44	...
4	Not Verified	Mar-2019	low_risk	n	18.76	...
...	...	...	...	...	...	...
68812	Source Verified	Jan-2019	low_risk	n	9.60	...
68813	Not Verified	Jan-2019	low_risk	n	29.07	...
68814	Source Verified	Jan-2019	low_risk	n	14.86	...
68815	Verified	Jan-2019	low_risk	n	9.96	...
68816	Verified	Jan-2019	low_risk	n	7.02	...

	total_bc_limit	total_il_high_credit_limit	hardship_flag	\
0	2000.0	61987.0	N	
1	41200.0	49197.0	N	
2	7500.0	43144.0	N	
3	19700.0	76506.0	N	
4	27600.0	20000.0	N	
...	...	...	...	
68812	11300.0	5425.0	N	
68813	13500.0	62939.0	N	
68814	3600.0	18492.0	N	
68815	100800.0	78634.0	N	
68816	23000.0	63090.0	N	

	debt_settlement_flag	issue_date	next_pymnt_date	issue_month	\
0	N	2019-03-01	2019-05-01	3	
1	N	2019-03-01	2019-05-01	3	
2	N	2019-03-01	2019-05-01	3	
3	N	2019-03-01	2019-05-01	3	
4	N	2019-03-01	2019-05-01	3	
...	...	...	...	...	
68812	N	2019-01-01	2019-05-01	1	
68813	N	2019-01-01	2019-05-01	1	
68814	N	2019-01-01	2019-05-01	1	
68815	N	2019-01-01	2019-05-01	1	
68816	N	2019-01-01	2019-05-01	1	

	issue_year	next_pymnt_d_month	next_pymnt_d_year
0	2019	5	2019
1	2019	5	2019
2	2019	5	2019
3	2019	5	2019
4	2019	5	2019
...	...	...	...
68812	2019	5	2019
68813	2019	5	2019
68814	2019	5	2019
68815	2019	5	2019
68816	2019	5	2019

```
[68817 rows x 92 columns]
```

### 1.3 Split the Data into Training and Testing

```
[8]: # Create our features
```

```
X = df[['loan_amnt', 'int_rate', 'installment', 'home_ownership', 'annual_inc',
        'verification_status', 'issue_d', 'pymnt_plan', 'dti',
        'delinq_2yrs', 'inq_last_6mths', 'open_acc', 'pub_rec', 'revol_bal',
        'total_acc', 'initial_list_status', 'out_prncp', 'out_prncp_inv',
        'total_pymnt', 'total_pymnt_inv', 'total_rec_prncp', 'total_rec_int',
        'total_rec_late_fee', 'recoveries', 'collection_recovery_fee',
        'last_pymnt_amnt', 'next_pymnt_d', 'collections_12_mths_ex_med',
        'policy_code', 'application_type', 'acc_now_delinq', 'tot_coll_amt',
        'tot_cur_bal', 'open_acc_6m', 'open_act_il', 'open_il_12m',
        'open_il_24m', 'mths_since_rcnt_il', 'total_bal_il', 'il_util',
        'open_rv_12m', 'open_rv_24m', 'max_bal_bc', 'all_util',
        'total_rev_hi_lim', 'inq_fi', 'total_cu_tl', 'inq_last_12m',
        'acc_open_past_24mths', 'avg_cur_bal', 'bc_open_to_buy', 'bc_util',
        'chargeoff_within_12_mths', 'delinq_amnt', 'mo_sin_old_il_acct',
        'mo_sin_old_rev_tl_op', 'mo_sin_rcnt_rev_tl_op', 'mo_sin_rcnt_tl',
        'mort_acc', 'mths_since_recent_bc', 'mths_since_recent_inq',
        'num_accts_ever_120_pd', 'num_actv_bc_tl', 'num_actv_rev_tl',
        'num_bc_sats', 'num_bc_tl', 'num_il_tl', 'num_op_rev_tl',
        'num_rev_accts', 'num_rev_tl_bal_gt_0', 'num_sats', 'num_tl_120dpd_2m',
        'num_tl_30dpd', 'num_tl_90g_dpd_24m', 'num_tl_op_past_12m',
        'pct_tl_nvr_dlq', 'percent_bc_gt_75', 'pub_rec_bankruptcies',
        'tax_liens', 'tot_hi_cred_lim', 'total_bal_ex_mort', 'total_bc_limit',
        'total_il_high_credit_limit', 'hardship_flag', 'issue_month', 'issue_year',
        'next_pymnt_d_month', 'next_pymnt_d_year', 'debt_settlement_flag']]

# Create our target
y = df["loan_status"]
```

```
[9]: y.value_counts()
```

```
[9]: low_risk      68470
     high_risk      347
     Name: loan_status, dtype: int64
```

```
[15]: categorical =
```

```
    ↳ ["home_ownership", "verification_status", "pymnt_plan", "application_type",
      ↳
    ↳ ["hardship_flag", "initial_list_status", "issue_month", "issue_year", "next_pymnt_d_month", "next
      ↳ "debt_settlement_flag"]
```

```
[16]: categorical_df = pd.get_dummies(df[categorical], prefix='d_')
categorical_df
```

```
[16]:
```

	d__ANY	d__MORTGAGE	d__OWN	d__RENT	d__Not Verified \
0	0	0	0	1	0
1	0	1	0	0	0
2	0	1	0	0	0
3	0	0	0	1	0
4	0	1	0	0	1
...	...	...	...	...	...
68812	0	0	0	1	0
68813	0	0	0	1	1
68814	0	1	0	0	0
68815	0	1	0	0	0
68816	0	1	0	0	0

	d__Source Verified	d__Verified	d__n	d__Individual	d__Joint App \
0	1	0	1	1	0
1	0	1	1	1	0
2	0	1	1	1	0
3	0	1	1	1	0
4	0	0	1	1	0
...	...	...	...	...	...
68812	1	0	1	1	0
68813	0	0	1	1	0
68814	1	0	1	1	0
68815	0	1	1	1	0
68816	0	1	1	1	0

	...	d__f	d__w	d__1	d__2	d__3	d__2019	d__4	d__5	d__2019	d__N
0	...	0	1	0	0	1	1	0	1	1	1
1	...	0	1	0	0	1	1	0	1	1	1
2	...	0	1	0	0	1	1	0	1	1	1
3	...	0	1	0	0	1	1	0	1	1	1
4	...	0	1	0	0	1	1	0	1	1	1
...	...	...	...	...	...	...	...	...	...	...	...
68812	...	0	1	1	0	0	1	0	1	1	1
68813	...	0	1	1	0	0	1	0	1	1	1
68814	...	0	1	1	0	0	1	0	1	1	1
68815	...	1	0	1	0	0	1	0	1	1	1
68816	...	0	1	1	0	0	1	0	1	1	1

[68817 rows x 21 columns]

```
[17]: print(X.shape)
```

(68817, 89)

```
[18]: X_new = pd.concat([X, categorical_df], axis="columns")
X_new.drop(categorical, axis="columns", inplace=True)
X_new.drop('issue_d', axis="columns", inplace=True)
X_new.drop('next_pymnt_d', axis="columns", inplace=True)
X = X_new
```

```
[19]: print(X.shape)
```

```
(68817, 97)
```

```
[20]: X.describe()
```

```
[20]:
```

	loan_amnt	int_rate	installment	annual_inc	dti \
count	68817.000000	68817.000000	68817.000000	6.881700e+04	68817.000000
mean	16677.594562	0.127718	480.652863	8.821371e+04	21.778153
std	10277.348590	0.048130	288.062432	1.155800e+05	20.199244
min	1000.000000	0.060000	30.890000	4.000000e+01	0.000000
25%	9000.000000	0.088100	265.730000	5.000000e+04	13.890000
50%	15000.000000	0.118000	404.560000	7.300000e+04	19.760000
75%	24000.000000	0.155700	648.100000	1.040000e+05	26.660000
max	40000.000000	0.308400	1676.230000	8.797500e+06	999.000000

	delinq_2yrs	inq_last_6mths	open_acc	pub_rec \
count	68817.000000	68817.000000	68817.000000	68817.000000
mean	0.217766	0.497697	12.587340	0.126030
std	0.718367	0.758122	6.022869	0.336797
min	0.000000	0.000000	2.000000	0.000000
25%	0.000000	0.000000	8.000000	0.000000
50%	0.000000	0.000000	11.000000	0.000000
75%	0.000000	1.000000	16.000000	0.000000
max	18.000000	5.000000	72.000000	4.000000

	revol_bal ...	d__f	d__w	d__1 \
count	68817.000000 ...	68817.000000	68817.000000	68817.000000
mean	17604.142828 ...	0.123879	0.876121	0.451066
std	21835.880400 ...	0.329446	0.329446	0.497603
min	0.000000 ...	0.000000	0.000000	0.000000
25%	6293.000000 ...	0.000000	1.000000	0.000000
50%	12068.000000 ...	0.000000	1.000000	0.000000
75%	21735.000000 ...	0.000000	1.000000	1.000000
max	587191.000000 ...	1.000000	1.000000	1.000000

	d__2	d__3	d__2019	d__4	d__5 \
count	68817.000000	68817.000000	68817.0	68817.000000	68817.000000
mean	0.371696	0.177238	1.0	0.383161	0.616839
std	0.483261	0.381873	0.0	0.486161	0.486161
min	0.000000	0.000000	1.0	0.000000	0.000000

25%	0.000000	0.000000	1.0	0.000000	0.000000
50%	0.000000	0.000000	1.0	0.000000	1.000000
75%	1.000000	0.000000	1.0	1.000000	1.000000
max	1.000000	1.000000	1.0	1.000000	1.000000

	d__2019	d__N
count	68817.0	68817.0
mean	1.0	1.0
std	0.0	0.0
min	1.0	1.0
25%	1.0	1.0
50%	1.0	1.0
75%	1.0	1.0
max	1.0	1.0

[8 rows x 97 columns]

```
[21]: # Check the balance of our target values
print(len(X))
print(len(y))
```

```
68817
68817
```

```
[22]: X
```

```
[22]:
```

	loan_amnt	int_rate	installment	annual_inc	dti	delinq_2yrs	\
0	10500.0	0.1719	375.35	66000.0	27.24	0.0	
1	25000.0	0.2000	929.09	105000.0	20.23	0.0	
2	20000.0	0.2000	529.88	56000.0	24.26	0.0	
3	10000.0	0.1640	353.55	92000.0	31.44	0.0	
4	22000.0	0.1474	520.39	52000.0	18.76	0.0	
...	...	...	...	...	...	...	
68812	10000.0	0.1502	346.76	26000.0	9.60	0.0	
68813	12000.0	0.2727	368.37	63000.0	29.07	0.0	
68814	5000.0	0.1992	185.62	52000.0	14.86	0.0	
68815	40000.0	0.0646	1225.24	520000.0	9.96	0.0	
68816	16000.0	0.1131	350.36	72000.0	7.02	2.0	

	inq_last_6mths	open_acc	pub_rec	revol_bal	...	d__f	d__w	d__1	\
0	0.0	8.0	0.0	1609.0	...	0	1	0	
1	0.0	17.0	1.0	18368.0	...	0	1	0	
2	0.0	8.0	0.0	13247.0	...	0	1	0	
3	1.0	10.0	1.0	17996.0	...	0	1	0	
4	1.0	14.0	0.0	9091.0	...	0	1	0	
...	...	...	...	...	...	...	...	...	
68812	0.0	9.0	0.0	2684.0	...	0	1	1	



68813	0.0	8.0	0.0	13314.0	...	0	1	1
68814	0.0	5.0	1.0	3715.0	...	0	1	1
68815	1.0	21.0	0.0	59529.0	...	1	0	1
68816	0.0	12.0	1.0	11882.0	...	0	1	1

	d__2	d__3	d__2019	d__4	d__5	d__2019	d__N
0	0	1	1	0	1	1	1
1	0	1	1	0	1	1	1
2	0	1	1	0	1	1	1
3	0	1	1	0	1	1	1
4	0	1	1	0	1	1	1
...	...	...	...	...	...	...	...
68812	0	0	1	0	1	1	1
68813	0	0	1	0	1	1	1
68814	0	0	1	0	1	1	1
68815	0	0	1	0	1	1	1
68816	0	0	1	0	1	1	1

[68817 rows x 97 columns]

```
[23]: X.to_csv("view.csv")
```

```
[24]: # Split the X and y into X_train, X_test, y_train, y_test
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=32)
```

## 1.4 Data Pre-Processing

Scale the training and testing data using the `StandardScaler` from `sklearn`. Remember that when scaling the data, you only scale the features data (`X_train` and `X_testing`).

```
[25]: # Create the StandardScaler instance
scaler = StandardScaler()
```

```
[26]: X_train.head()
```

```
[26]:      loan_amnt  int_rate  installment  annual_inc  dti  delinq_2yrs  \
63152    15000.0    0.2235      417.28    125000.0  11.06         0.0
2905     40000.0    0.1033      856.40    105000.0  15.02         0.0
27484    12925.0    0.1691      460.24     55000.0  13.81         0.0
12354    10000.0    0.1474      345.39     70000.0   9.19         0.0
4134     15000.0    0.1171      496.14     85000.0  10.29         0.0

      inq_last_6mths  open_acc  pub_rec  revol_bal  ...  d__f  d__w  d__1  \
63152              2.0      10.0       2.0     9550.0  ...    0    1    1
2905              1.0      13.0       0.0    34395.0  ...    0    1    0
27484              2.0       9.0       0.0     8524.0  ...    0    1    0
```

12354	1.0	9.0	1.0	2352.0	...	1	0	0
4134	0.0	8.0	0.0	1701.0	...	0	1	0

	d__2	d__3	d__2019	d__4	d__5	d__2019	d__N
63152	0	0	1	0	1	1	1
2905	0	1	1	0	1	1	1
27484	1	0	1	0	1	1	1
12354	1	0	1	0	1	1	1
4134	0	1	1	0	1	1	1

[5 rows x 97 columns]

```
[27]: # Fit the Standard Scaler with the training data
# When fitting scaling functions, only train on the training dataset
X_scaler = scaler.fit(X_train)
# scaler_test = StandardScaler().fit(X_test)
```

```
[28]: # Scale the training and testing data
X_train_scaled = X_scaler.transform(X_train)
X_test_scaled = X_scaler.transform(X_test)
```

## 1.5 Ensemble Learners

In this section, you will compare two ensemble algorithms to determine which algorithm results in the best performance. You will train a Balanced Random Forest Classifier and an Easy Ensemble classifier. For each algorithm, be sure to complete the following steps:

1. Train the model using the training data.
2. Calculate the balanced accuracy score from `sklearn.metrics`.
3. Display the confusion matrix from `sklearn.metrics`.
4. Generate a classification report using the `imbalanced_classification_report` from `imbalanced-learn`.
5. For the Balanced Random Forest Classifier only, print the feature importance sorted in descending order (most important feature to least important) along with the feature score

Note: Use a random state of 1 for each algorithm to ensure consistency between tests

### 1.5.1 Balanced Random Forest Classifier

```
[35]: # Resample the training data with the BalancedRandomForestClassifier
brfc = BalancedRandomForestClassifier(max_depth=2, random_state=1)
brfc.fit(X_train, y_train)
```

```
[35]: BalancedRandomForestClassifier(max_depth=2, random_state=1)
```

```
[36]: # Calculated the balanced accuracy score
y_pred = brfc.predict(X_test)
balanced_accuracy_score(y_test, y_pred)
```

```
[36]: 0.7147418320530224
```

```
[37]: # Display the confusion matrix
confusion_matrix(y_test, y_pred)
```


```
[37]: array([[ 38,   44],
          [ 581, 16542]])
```

```
[38]: # Print the imbalanced classification report
print(classification_report_imbalanced(y_test, y_pred))
```

	pre	rec	spe	f1	geo	iba
sup						
high_risk	0.06	0.46	0.97	0.11	0.67	0.43
low_risk	1.00	0.97	0.46	0.98	0.67	0.47
avg / total	0.99	0.96	0.47	0.98	0.67	0.47

```
[39]: # List the features sorted in descending order by feature importance
importances_df = pd.DataFrame(sorted(zip(clf.feature_importances_, X.columns),
reverse=True))
importances_df.set_index(importances_df[1], inplace=True)
importances_df.drop(columns=1, inplace=True)
importances_df.rename(columns={0: 'Feature Importances'}, inplace=True)
importances_sorted = importances_df.sort_values(by='Feature Importances',
ascending=True)
importances_sorted.plot(kind='barh', color='lightgreen', title= 'Features
Importances', legend=False, figsize=(14,20))
```

```
[39]: <AxesSubplot:title={'center': 'Features Importances'}, ylabel='1'>
```



## 1.5.2 Easy Ensemble Classifier

```
[46]: # Train the Classifier
eec = EasyEnsembleClassifier(random_state=1)
eec.fit(X_train, y_train)
```

```
[46]: EasyEnsembleClassifier(random_state=1)
```

```
[47]: # Calculated the balanced accuracy score
y_pred = eec.predict(X_test)
balanced_accuracy_score(y_test, y_pred)
```

```
[47]: 0.9197705838531258
```

```
[48]: # Display the confusion matrix
confusion_matrix(y_test, y_pred)
```

```
[48]: array([[ 74,    8],
       [1077, 16046]])
```

```
[49]: # Print the imbalanced classification report
print(classification_report_imbalanced(y_test, y_pred))
```

	pre	rec	spe	f1	geo	iba
sup						
high_risk	0.06	0.90	0.94	0.12	0.92	0.84
82						
low_risk	1.00	0.94	0.90	0.97	0.92	0.85
17123						
avg / total	1.00	0.94	0.90	0.96	0.92	0.85
17205						

```
[52]: print(np.mean([est.steps[1][1].feature_importances_ for est in eec.
    ↳ estimators_], axis=0))
```

```
[0.014 0.068 0.076 0.    0.008 0.002 0.006 0.002 0.    0.002 0.002 0.016
0.014 0.022 0.026 0.096 0.168 0.018 0.    0.    0.106 0.    0.    0.
0.004 0.004 0.004 0.002 0.    0.    0.008 0.004 0.    0.    0.    0.012
0.006 0.002 0.004 0.    0.    0.002 0.006 0.016 0.016 0.    0.    0.002
0.01  0.    0.004 0.002 0.006 0.002 0.    0.    0.    0.    0.002 0.002
0.002 0.    0.002 0.004 0.    0.    0.    0.004 0.    0.002 0.    0.
0.004 0.004 0.008 0.    0.    0.    0.    0.    0.002 0.    0.    0.
0.    0.    0.    0.    0.    0.078 0.004 0.048 0.    0.034 0.038 0.
0.    ]
```

### 1.5.3 Final Questions

1. Which model had the best balanced accuracy score?

Easy Ensemble Classifier with 0.9197

2. Which model had the best recall score?

Balanced Random Forest Classifier with 0.96 (96%)

3. Which model had the best geometric mean score?

Easy Ensemble Classifier with 0.92

4. What are the top three features?

last\_paymnt\_amt, total\_payment, total\_rec\_int

[ ]: